



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

Data Operations and Parallel Mapping

Parallel Programming in Scala

Viktor Kuncak

Parallelism and collections

Parallel processing of collections is important

- ▶ one the main applications of parallelism today

We examine conditions when this can be done

- ▶ properties of collections: ability to split, combine
- ▶ properties of operations: associativity, independence

Functional programming and collections

Operations on collections are key to functional programming

map: apply function to each element

- ▶ `List(1,3,8).map(x => x*x) == List(1, 9, 64)`

Functional programming and collections

Operations on collections are key to functional programming

map: apply function to each element

- ▶ `List(1,3,8).map(x => x*x) == List(1, 9, 64)`

fold: combine elements with a given operation

- ▶ `List(1,3,8).fold(100)((s,x) => s + x) == 112`

Functional programming and collections

Operations on collections are key to functional programming

map: apply function to each element

- ▶ `List(1,3,8).map(x => x*x) == List(1, 9, 64)`

fold: combine elements with a given operation

- ▶ `List(1,3,8).fold(100)((s,x) => s + x) == 112`

scan: combine folds of all list prefixes

- ▶ `List(1,3,8).scan(100)((s,x) => s + x) == List(100, 101, 104, 112)`

These operations are even more important for parallel than sequential collections: they encapsulate more complex algorithms

Choice of data structures

We use **List** to specify the results of operations

Lists are not good for parallel implementations because we cannot efficiently

- ▶ split them in half (need to search for the middle)
- ▶ combine them (concatenation needs linear time)

Choice of data structures

We use **List** to specify the results of operations

Lists are not good for parallel implementations because we cannot efficiently

- ▶ split them in half (need to search for the middle)
- ▶ combine them (concatenation needs linear time)

We use for now these alternatives

- ▶ **arrays**: imperative (recall array sum)
- ▶ **trees**: can be implemented functionally

Subsequent lectures examine Scala's parallel collection libraries

- ▶ includes many more data structures, implemented efficiently

Map: meaning and properties

Map applies a given function to each list element

`List(1,3,8).map(x => x*x) == List(1, 9, 64)`

$\text{List}(a_1, a_2, \dots, a_n).\text{map}(f) == \text{List}(f(a_1), f(a_2), \dots, f(a_n))$

Properties to keep in mind:

- ▶ `list.map(x => x) == list`
- ▶ `list.map(f.compose(g)) == list.map(g).map(f)`

Recall that $(f.\text{compose}(g))(x) = f(g(x))$

Map as function on lists

Sequential definition:


```
def mapSeq[A,B](lst: List[A], f : A => B): List[B] = lst match {  
  case Nil => Nil  
  case h :: t => f(h) :: mapSeq(t,f)  
}
```

We would like a version that parallelizes

- ▶ computations of $f(h)$ for different elements h
- ▶ finding the elements themselves (list is not a good choice)

Sequential map of an array producing an array

```
def mapASegSeq[A,B](inp: Array[A], left: Int, right: Int, f : A => B,  
                    out: Array[B]) = {  
  // Writes to out(i) for left <= i <= right-1  
  var i= left  
  while (i < right) {  
    out(i)= f(inp(i))  
    i= i+1  
  } }  
val in= Array(2,3,4,5,6)  
val out= Array(0,0,0,0,0)  
val f= (x:Int) => x*x  
mapASegSeq(in, 1, 3, f, out)  
out
```




```
res1: Array[Int] = Array(0, 9, 16, 0, 0)
```

Parallel map of an array producing an array

```
def mapASegPar[A,B](inp: Array[A], left: Int, right: Int, f : A => B,
                    out: Array[B]): Unit = {
  // Writes to out(i) for left <= i <= right-1
  if (right - left < threshold)
    mapASegSeq(inp, left, right, f, out)
  else {
    val mid = left + (right - left)/2
    parallel(mapASegPar(inp, left, mid, f, out),
             mapASegPar(inp, mid, right, f, out))
  }
}
```

Note:



- ▶ writes need to be disjoint (otherwise: non-deterministic behavior)
- ▶ threshold needs to be large enough (otherwise we lose efficiency)

Example of using mapASegPar: pointwise exponent

Raise each array element to power p :

$$\text{Array}(a_1, a_2, \dots, a_n) \longrightarrow \text{Array}(|a_1|^p, |a_2|^p, \dots, |a_n|^p)$$

We can use previously defined higher-order functions:

```
val p: Double = 1.5
```

```
def f(x: Int): Double = power(x, p)
```

```
mapASegSeq(inp, 0, inp.length, f, out)    // sequential
```

```
mapASegPar(inp, 0, inp.length, f, out)    // parallel
```

Example of using mapASegPar: pointwise exponent

Raise each array element to power p :

$$\text{Array}(a_1, a_2, \dots, a_n) \longrightarrow \text{Array}(|a_1|^p, |a_2|^p, \dots, |a_n|^p)$$

We can use previously defined higher-order functions:

```
val p: Double = 1.5
```

```
def f(x: Int): Double = power(x, p)
```

```
mapASegSeq(inp, 0, inp.length, f, out)    // sequential
```

```
mapASegPar(inp, 0, inp.length, f, out)    // parallel
```

Questions on performance:

- ▶ are there performance gains from parallel execution
- ▶ performance of re-using higher-order functions vs re-implementing

Sequential pointwise exponent written from scratch

```
def normsOf(inp: Array[Int], p: Double,  
            left: Int, right: Int,  
            out: Array[Double]): Unit = {  
  var i = left  
  while (i < right) {  
    out(i) = power(inp(i), p)  
    i = i + 1  
  }  
}
```

Parallel pointwise exponent written from scratch

```
def normsOfPar(inp: Array[Int], p: Double,
               left: Int, right: Int,
               out: Array[Double]): Unit = {
  if (right - left < threshold) {
    var i = left
    while (i < right) {
      out(i) = power(inp(i), p)
      i = i + 1
    }
  } else {
    val mid = left + (right - left) / 2
    parallel(normsOfPar(inp, p, left, mid, out),
             normsOfPar(inp, p, mid, right, out))
  }
}
```

Measured performance using scalameter

- ▶ `inp.length = 2000000`
- ▶ `threshold = 10000`
- ▶ Intel(R) Core(TM) i7-3770K CPU @ 3.50GHz (4-core, 8 HW threads), 16GB RAM

<i>expression</i>	<i>time(ms)</i>
<i>mapASegSeq(inp, 0, inp.length, f, out)</i>	174.17
<i>mapASegPar(inp, 0, inp.length, f, out)</i>	28.93
<i>normsOfSeq(inp, p, 0, inp.length, out)</i>	166.84
<i>normsOfPar(inp, p, 0, inp.length, out)</i>	28.17

Measured performance using scalameter

- ▶ `inp.length = 2000000`
- ▶ `threshold = 10000`
- ▶ Intel(R) Core(TM) i7-3770K CPU @ 3.50GHz (4-core, 8 HW threads), 16GB RAM

<i>expression</i>	<i>time(ms)</i>
<i>mapASegSeq(inp, 0, inp.length, f, out)</i>	174.17
<i>mapASegPar(inp, 0, inp.length, f, out)</i>	28.93
<i>normsOfSeq(inp, p, 0, inp.length, out)</i>	166.84
<i>normsOfPar(inp, p, 0, inp.length, out)</i>	28.17

Parallelization pays off

Manually removing higher-order functions does not pay off

Parallel map on immutable trees

Consider trees where

- ▶ leaves store array segments
- ▶ non-leaf node stores two subtrees

```
sealed abstract class Tree[A] { val size: Int }  
case class Leaf[A](a: Array[A]) extends Tree[A] {  
  override val size = a.size  
}  
case class Node[A](l: Tree[A], r: Tree[A]) extends Tree[A] {  
  override val size = l.size + r.size  
}
```

Assume that our trees are balanced: we can explore branches in parallel

Parallel map on immutable trees

```
def mapTreePar[A:Manifest,B:Manifest](t: Tree[A], f: A => B) : Tree[B] =  
t match {  
  case Leaf(a) => {  
    val len = a.length; val b = new Array[B](len)  
    var i= 0  
    while (i < len) { b(i)= f(a(i)); i= i + 1 }  
    Leaf(b) }  
  case Node(l,r) => {  
    val (lb,rb) = parallel(mapTreePar(l,f), mapTreePar(r,f))  
    Node(lb, rb) }  
}
```

Speedup and performance similar as for the array

Give depth bound of mapTreePar

Give a correct but as tight as possible asymptotic parallel computation depth bound for `mapTreePar` applied to complete trees with height h and 2^h nodes, assuming the passed first-class function f executes in constant time.

1. 2^h
2. h
3. $\log h$
4. $h \log h$
5. $h2^h$

Give depth bound of mapTreePar

Give a correct but as tight as possible asymptotic parallel computation depth bound for `mapTreePar` applied to complete trees with height h and 2^h nodes, assuming the passed first-class function f executes in constant time.

1. 2^h
2. h
3. $\log h$
4. $h \log h$
5. $h2^h$

Answer: h . The computation depth equals the height of the tree.

Comparison of arrays and immutable trees

Arrays:

- ▶ (+) random access to elements, on shared memory can share array
- ▶ (+) good memory locality
- ▶ (-) imperative: must ensure parallel tasks write to disjoint parts
- ▶ (-) expensive to concatenate

Immutable trees:

- ▶ (+) purely functional, produce new trees, keep old ones
- ▶ (+) no need to worry about disjointness of writes by parallel tasks
- ▶ (+) efficient to combine two trees
- ▶ (-) high memory allocation overhead
- ▶ (-) bad locality

End of Slide Deck