



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

Implementing Combiners

Parallel Programming in Scala

Aleksandar Prokopec

Builders

Builders are used in sequential collection methods:

Builders

Builders are used in sequential collection methods:

```
trait Builder[T, Repr] {  
  def +=(elem: T): this.type  
  def result: Repr  
}
```

Combiners

```
trait Combiner[T, Repr] extends Builder[T, Repr] {  
  def combine(that: Combiner[T, Repr]): Combiner[T, Repr]  
}
```

Combiners

```
trait Combiner[T, Repr] extends Builder[T, Repr] {  
  def combine(that: Combiner[T, Repr]): Combiner[T, Repr]  
}
```

How can we implement the combine method *efficiently*?

Combiners

- ▶ when Repr is a set or a map, combine represents union

Combiners

- ▶ when Repr is a set or a map, combine represents union
- ▶ when Repr is a sequence, combine represents concatenation

Combiners

The combine operation must be efficient, i.e. execute in $O(\log n + \log m)$ time, where n and m are the sizes of two input combiners.

Combiners

The combine operation must be efficient, i.e. execute in $O(\log n + \log m)$ time, where n and m are the sizes of two input combiners.

Question: Is the method combine *efficient*?

```
def combine(xs: Array[Int], ys: Array[Int]): Array[Int] = {  
  val r = new Array[Int](xs.length + ys.length)  
  Array.copy(xs, 0, r, 0, xs.length)  
  Array.copy(ys, 0, r, xs.length, ys.length)  
  r  
}
```

- ▶ Yes.
- ▶ No.

Array Concatenation

Arrays cannot be efficiently concatenated.

Sets

Typically, set data structures have efficient lookup, insertion and deletion.

Sets

Typically, set data structures have efficient lookup, insertion and deletion.

- ▶ hash tables – expected $O(1)$

Sets

Typically, set data structures have efficient lookup, insertion and deletion.

- ▶ hash tables – expected $O(1)$
- ▶ balanced trees – $O(\log n)$

Sets

Typically, set data structures have efficient lookup, insertion and deletion.

- ▶ hash tables – expected $O(1)$
- ▶ balanced trees – $O(\log n)$
- ▶ linked lists – $O(n)$

Sets

Typically, set data structures have efficient lookup, insertion and deletion.

- ▶ hash tables – expected $O(1)$
- ▶ balanced trees – $O(\log n)$
- ▶ linked lists – $O(n)$

Most set implementations do not have efficient union operation.

Sequences

Operation complexity for sequences can vary.

Sequences

Operation complexity for sequences can vary.

- ▶ mutable linked lists – $O(1)$ prepend and append, $O(n)$ insertion

Sequences

Operation complexity for sequences can vary.

- ▶ mutable linked lists – $O(1)$ prepend and append, $O(n)$ insertion
- ▶ functional (cons) lists – $O(1)$ prepend operations, everything else $O(n)$

Sequences

Operation complexity for sequences can vary.

- ▶ mutable linked lists – $O(1)$ prepend and append, $O(n)$ insertion
- ▶ functional (cons) lists – $O(1)$ prepend operations, everything else $O(n)$
- ▶ array lists – amortized $O(1)$ append, $O(1)$ random access, otherwise $O(n)$

Sequences

Operation complexity for sequences can vary.

- ▶ mutable linked lists – $O(1)$ prepend and append, $O(n)$ insertion
- ▶ functional (cons) lists – $O(1)$ prepend operations, everything else $O(n)$
- ▶ array lists – amortized $O(1)$ append, $O(1)$ random access, otherwise $O(n)$

Mutable linked list can have $O(1)$ concatenation, but for most sequences, concatenation is $O(n)$.